

```

import tensorflow as tf

from tensorflow.keras import datasets, layers, models

from tensorflow.keras.utils import to_categorical


# Load and preprocess the MNIST dataset
(train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()

train_images = train_images.reshape(-1, 28, 28, 1).astype('float32') / 255.0
test_images = test_images.reshape(-1, 28, 28, 1).astype('float32') / 255.0
train_labels = to_categorical(train_labels, 10)
test_labels = to_categorical(test_labels, 10)


# Build a simple CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])


# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',

```

```
metrics=['accuracy'])
```

```
# Train the model
```

```
model.fit(train_images, train_labels, epochs=5, batch_size=64, validation_split=0.1)
```

```
# Evaluate the model
```

```
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=0)
```

```
print(f'Test accuracy: {test_acc:.4f}')
```